

(a centralized global routing algorithm) or replicated at multiple sites. The key distinguishing feature here, however, is that a global algorithm has complete information about connectivity and link costs. In practice, algorithms with global state information are often referred to as **link-state (LS) algorithms**, since the algorithm must be aware of the cost of each link in the network. We'll study LS algorithms in Section 4.5.1.

- In a **decentralized routing algorithm**, the calculation of the least-cost path is carried out in an iterative, distributed manner. No node has complete information about the costs of all network links. Instead, each node begins with only the knowledge of the costs of its own directly attached links. Then, through an iterative process of calculation and exchange of information with its neighboring nodes (that is, nodes that are at the other end of links to which it itself is attached), a node gradually calculates the least-cost path to a destination or set of destinations. The decentralized routing algorithm we'll study below in Section 4.5.2 is called a distance-vector (DV) algorithm, because each node maintains a vector of estimates of the costs (distances) to all other nodes in the network.

A second broad way to classify routing algorithms is according to whether they are static or dynamic. In **static routing algorithms**, routes change very slowly over time, often as a result of human intervention (for example, a human manually editing a router's forwarding table). **Dynamic routing algorithms** change the routing paths as the network traffic loads or topology change. A dynamic algorithm can be run either periodically or in direct response to topology or link cost changes. While dynamic algorithms are more responsive to network changes, they are also more susceptible to problems such as routing loops and oscillation in routes.

A third way to classify routing algorithms is according to whether they are load-sensitive or load-insensitive. In a **load-sensitive algorithm**, link costs vary dynamically to reflect the current level of congestion in the underlying link. If a high cost is associated with a link that is currently congested, a routing algorithm will tend to choose routes around such a congested link. While early ARPAnet routing algorithms were load-sensitive [McQuillan 1980], a number of difficulties were encountered [Huijema 1998]. Today's Internet routing algorithms (such as RIP, OSPF, and BGP) are **load-insensitive**, as a link's cost does not explicitly reflect its current (or recent past) level of congestion.

### 4.5.1 The Link-State (LS) Routing Algorithm

Recall that in a link-state algorithm, the network topology and all link costs are known, that is, available as input to the LS algorithm. In practice this is accomplished by having each node broadcast link-state packets to *all* other nodes in the network, with each link-state packet containing the identities and costs of its attached links. In practice (for example, with the Internet's OSPF routing protocol, discussed in Section 4.6.1) this is often accomplished by a **link-state broadcast**.

algorithm [Perlman 1999]. We'll cover broadcast algorithms in Section 4.7. The result of the nodes' broadcast is that all nodes have an identical and complete view of the network. Each node can then run the LS algorithm and compute the same set of least-cost paths as every other node.

The link-state routing algorithm we present below is known as *Dijkstra's algorithm*, named after its inventor. A closely related algorithm is Prim's algorithm; see [Cormen 2001] for a general discussion of graph algorithms. Dijkstra's algorithm computes the least-cost path from one node (the source, which we will refer to as  $u$ ) to all other nodes in the network. Dijkstra's algorithm is iterative and has the property that after the  $k$ th iteration of the algorithm, the least-cost paths are known to  $k$  destination nodes, and among the least-cost paths to all destination nodes, these  $k$  paths will have the  $k$  smallest costs. Let us define the following notation:

- $D(v)$ : cost of the least-cost path from the source node to destination  $v$  as of this iteration of the algorithm.
- $p(v)$ : previous node (neighbor of  $v$ ) along the current least-cost path from the source to  $v$ .
- $N'$  : subset of nodes;  $v$  is in  $N'$  if the least-cost path from the source to  $v$  is definitively known.

The global routing algorithm consists of an initialization step followed by a loop. The number of times the loop is executed is equal to the number of nodes in the network. Upon termination, the algorithm will have calculated the shortest paths from the source node  $u$  to every other node in the network.

### Link-State (LS) Algorithm for Source Node $u$

```

1 Initialization:
2   N' = {u}
3   for all nodes v
4     if v is a neighbor of u
5       then D(v) = c(u,v)
6     else D(v) = ∞
7
8 Loop
9   find w not in N' such that D(w) is a minimum
10  add w to N'
11  update D(v) for each neighbor v of w and not in N':
12    D(v) = min( D(v), D(w) + c(w,v) )
13  /* new cost to v is either old cost to v or known
14    least path cost to w plus cost from w to v */
15 until N' = N

```



As an example, let's consider the network in Figure 4.27 and compute the least-cost paths from  $u$  to all possible destinations. A tabular summary of the algorithm's computation is shown in Table 4.3, where each line in the table gives the values of the algorithm's variables at the end of the iteration. Let's consider the few first steps in detail.

- In the initialization step, the currently known least-cost paths from  $u$  to its directly attached neighbors,  $v$ ,  $x$ , and  $w$ , are initialized to 2, 1, and 5, respectively. Note in particular that the cost to  $w$  is set to 5 (even though we will soon see that a lesser-cost path does indeed exist) since this is the cost of the direct (one hop) link from  $u$  to  $w$ . The costs to  $y$  and  $z$  are set to infinity because they are not directly connected to  $u$ .
- In the first iteration, we look among those nodes not yet added to the set  $N'$  and find that node with the least cost as of the end of the previous iteration. That node is  $x$ , with a cost of 1, and thus  $x$  is added to the set  $N'$ . Line 12 of the LS algorithm is then performed to update  $D(v)$  for all nodes  $v$ , yielding the results shown in the second line (Step 1) in Table 4.3. The cost of the path to  $v$  is unchanged. The cost of the path to  $w$  (which was 5 at the end of the initialization) through node  $x$  is found to have a cost of 4. Hence this lower-cost path is selected and  $w$ 's predecessor along the shortest path from  $u$  is set to  $x$ . Similarly, the cost to  $y$  (through  $x$ ) is computed to be 2, and the table is updated accordingly.
- In the second iteration, nodes  $v$  and  $y$  are found to have the least-cost paths (2), and we break the tie arbitrarily and add  $y$  to the set  $N'$  so that  $N'$  now contains  $u$ ,  $x$ , and  $y$ . The cost to the remaining nodes not yet in  $N'$ , that is, nodes  $v$ ,  $w$ , and  $z$ , are updated via line 12 of the LS algorithm, yielding the results shown in the third row in the Table 4.3.
- And so on. . . .

When the LS algorithm terminates, we have, for each node, its predecessor along the least-cost path from the source node. For each predecessor, we also

step	$N'$	$D(v), p(v)$	$D(w), p(w)$	$D(x), p(x)$	$D(y), p(y)$	$D(z), p(z)$
0	$u$	2, $u$	5, $u$	1, $u$	$\infty$	$\infty$
1	$ux$	2, $u$	4, $x$		2, $x$	$\infty$
2	$uxy$	2, $u$	3, $y$			4, $y$
3	$uxyv$		3, $y$			4, $y$
4	$uxyvw$					4, $y$
5	$uxyvwz$					

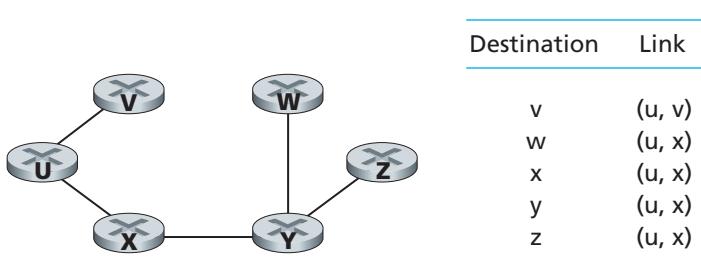
**Table 4.3** ♦ Running the link-state algorithm on the network in Figure 4.27

have *its* predecessor, and so in this manner we can construct the entire path from the source to all destinations. The forwarding table in a node, say node  $u$ , can then be constructed from this information by storing, for each destination, the next-hop node on the least-cost path from  $u$  to the destination. Figure 4.28 shows the resulting least-cost paths and forwarding table in  $u$  for the network in Figure 4.27.

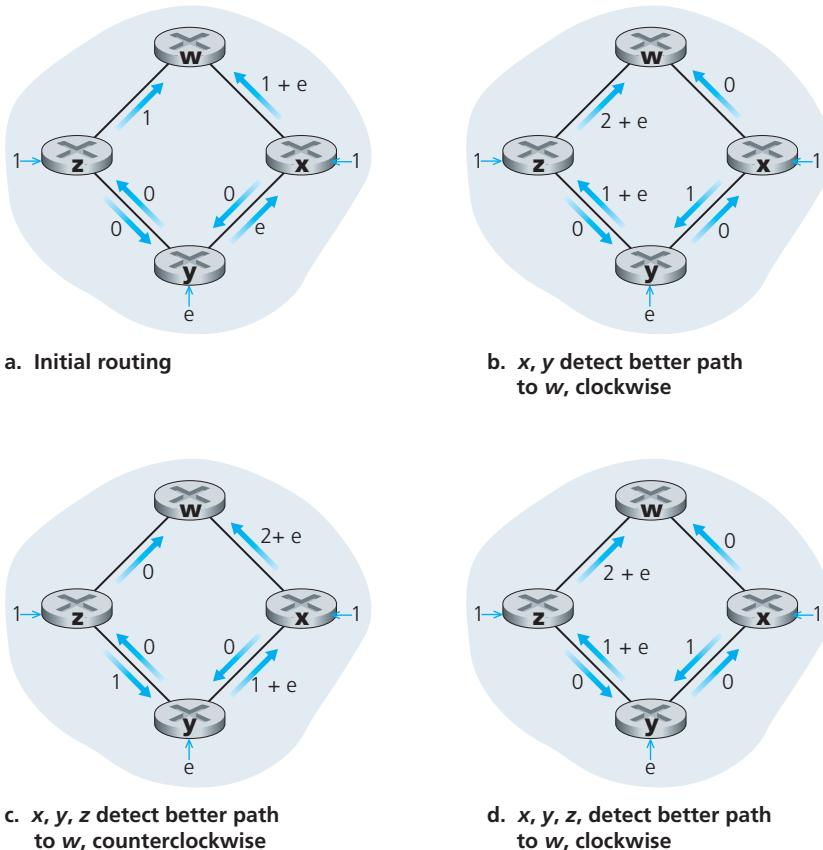
What is the computational complexity of this algorithm? That is, given  $n$  nodes (not counting the source), how much computation must be done in the worst case to find the least-cost paths from the source to all destinations? In the first iteration, we need to search through all  $n$  nodes to determine the node,  $w$ , not in  $N'$  that has the minimum cost. In the second iteration, we need to check  $n - 1$  nodes to determine the minimum cost; in the third iteration  $n - 2$  nodes, and so on. Overall, the total number of nodes we need to search through over all the iterations is  $n(n + 1)/2$ , and thus we say that the preceding implementation of the LS algorithm has worst-case complexity of order  $n$  squared:  $O(n^2)$ . (A more sophisticated implementation of this algorithm, using a data structure known as a heap, can find the minimum in line 9 in logarithmic rather than linear time, thus reducing the complexity.)

Before completing our discussion of the LS algorithm, let us consider a pathology that can arise. Figure 4.29 shows a simple network topology where link costs are equal to the load carried on the link, for example, reflecting the delay that would be experienced. In this example, link costs are not symmetric; that is,  $c(u,v)$  equals  $c(v,u)$  only if the load carried on both directions on the link  $(u,v)$  is the same. In this example, node  $z$  originates a unit of traffic destined for  $w$ , node  $x$  also originates a unit of traffic destined for  $w$ , and node  $y$  injects an amount of traffic equal to  $e$ , also destined for  $w$ . The initial routing is shown in Figure 4.29(a) with the link costs corresponding to the amount of traffic carried.

When the LS algorithm is next run, node  $y$  determines (based on the link costs shown in Figure 4.29(a)) that the clockwise path to  $w$  has a cost of 1, while the counterclockwise path to  $w$  (which it had been using) has a cost of  $1 + e$ . Hence  $y$ 's



**Figure 4.28** ♦ Least cost path and forwarding table for node  $u$



**Figure 4.29** ♦ Oscillations with congestion-sensitive routing

least-cost path to  $w$  is now clockwise. Similarly,  $x$  determines that its new least-cost path to  $w$  is also clockwise, resulting in costs shown in Figure 4.29(b). When the LS algorithm is run next, nodes  $x$ ,  $y$ , and  $z$  all detect a zero-cost path to  $w$  in the counterclockwise direction, and all route their traffic to the counterclockwise routes. The next time the LS algorithm is run,  $x$ ,  $y$ , and  $z$  all then route their traffic to the clockwise routes.

What can be done to prevent such oscillations (which can occur in any algorithm, not just an LS algorithm, that uses a congestion or delay-based link metric)? One solution would be to mandate that link costs not depend on the amount of traffic carried—an unacceptable solution since one goal of routing is to avoid