

check used at the link layer? Recall that the transport layer is typically implemented in software in a host as part of the host's operating system. Because transport-layer error detection is implemented in software, it is important to have a simple and fast error-detection scheme such as checksumming. On the other hand, error detection at the link layer is implemented in dedicated hardware in adapters, which can rapidly perform the more complex CRC operations. Feldmeier [Feldmeier 1995] presents fast software implementation techniques for not only weighted checksum codes, but CRC (see below) and other codes as well.


### 5.2.3 Cyclic Redundancy Check (CRC)

An error-detection technique used widely in today's computer networks is based on **cyclic redundancy check (CRC) codes**. CRC codes are also known as **polynomial codes**, since it is possible to view the bit string to be sent as a polynomial whose coefficients are the 0 and 1 values in the bit string, with operations on the bit string interpreted as polynomial arithmetic.

CRC codes operate as follows. Consider the  $d$ -bit piece of data,  $D$ , that the sending node wants to send to the receiving node. The sender and receiver must first agree on an  $r + 1$  bit pattern, known as a **generator**, which we will denote as  $G$ . We will require that the most significant (leftmost) bit of  $G$  be a 1. The key idea behind CRC codes is shown in Figure 5.6. For a given piece of data,  $D$ , the sender will choose  $r$  additional bits,  $R$ , and append them to  $D$  such that the resulting  $d + r$  bit pattern (interpreted as a binary number) is exactly divisible by  $G$  (i.e., has no remainder) using modulo-2 arithmetic. The process of error checking with CRCs is thus simple: The receiver divides the  $d + r$  received bits by  $G$ . If the remainder is nonzero, the receiver knows that an error has occurred; otherwise the data is accepted as being correct.

All CRC calculations are done in modulo-2 arithmetic without carries in addition or borrows in subtraction. This means that addition and subtraction are identical, and both are equivalent to the bitwise exclusive-or (XOR) of the operands. Thus, for example,

$$\begin{array}{rcl} 1011 & \text{XOR} & 0101 = 1110 \\ 1001 & \text{XOR} & 1101 = 0100 \end{array}$$



**Figure 5.6** ◆ CRC

Also, we similarly have

$$\begin{array}{r} 1011 - 0101 = 1110 \\ 1001 - 1101 = 0100 \end{array}$$

Multiplication and division are the same as in base-2 arithmetic, except that any required addition or subtraction is done without carries or borrows. As in regular binary arithmetic, multiplication by  $2^k$  left shifts a bit pattern by  $k$  places. Thus, given  $D$  and  $R$ , the quantity  $D \cdot 2^r$  XOR  $R$  yields the  $d + r$  bit pattern shown in Figure 5.6. We'll use this algebraic characterization of the  $d + r$  bit pattern from Figure 5.6 in our discussion below.

Let us now turn to the crucial question of how the sender computes  $R$ . Recall that we want to find  $R$  such that there is an  $n$  such that

$$D \cdot 2^r \text{ XOR } R = nG$$

That is, we want to choose  $R$  such that  $G$  divides into  $D \cdot 2^r$  XOR  $R$  without remainder. If we XOR (that is, add modulo-2, without carry)  $R$  to both sides of the above equation, we get

$$D \cdot 2^r = nG \text{ XOR } R$$

This equation tells us that if we divide  $D \cdot 2^r$  by  $G$ , the value of the remainder is precisely  $R$ . In other words, we can calculate  $R$  as


$$R = \text{remainder} \frac{D \cdot 2^r}{G}$$

Figure 5.7 illustrates this calculation for the case of  $D = 101110$ ,  $d = 6$ ,  $G = 1001$ , and  $r = 3$ . The 9 bits transmitted in this case are 101110 011. You should check these calculations for yourself and also check that indeed  $D \cdot 2^r = 101011 \cdot G$  XOR  $R$ .

International standards have been defined for 8-, 12-, 16-, and 32-bit generators,  $G$ . The CRC-32 32-bit standard, which has been adopted in a number of link-level IEEE protocols, uses a generator of

$$G_{\text{CRC-32}} = 100000100110000010001110110110111$$

Each of the CRC standards can detect burst errors of fewer than  $r + 1$  bits. (This means that all consecutive bit errors of  $r$  bits or fewer will be detected.) Furthermore, under appropriate assumptions, a burst of length greater than  $r + 1$  bits is detected with probability  $1 - 0.5^r$ . Also, each of the CRC standards can detect any odd number of bit errors. See [Williams 1993] for a discussion of implementing CRC checks. The theory



**Figure 5.7** ♦ A sample CRC calculation

behind CRC codes and even more powerful codes is beyond the scope of this text. The text [Schwartz 1980] provides an excellent introduction to this topic.

## 5.3 Multiple Access Links and Protocols

In the introduction to this chapter, we noted that there are two types of network links: point-to-point links and broadcast links. A **point-to-point link** consists of a single sender at one end of the link and a single receiver at the other end of the link. Many link-layer protocols have been designed for point-to-point links; the point-to-point protocol (PPP) and high-level data link control (HDLC) are two such protocols that we'll cover later in this chapter. The second type of link, a **broadcast link**, can have multiple sending and receiving nodes all connected to the same, single, shared broadcast channel. The term *broadcast* is used here because when any one node transmits a frame, the channel broadcasts the frame and each of the other nodes receives a copy. Ethernet and wireless LANs are examples of broadcast link-layer technologies. In this section we'll take a step back from specific link-layer protocols and first examine a problem of central importance to the link layer: how to coordinate the access of multiple sending and receiving nodes to a shared broadcast channel—the **multiple access problem**. Broadcast channels are often used in LANs, networks that are geographically concentrated in a single building (or on a corporate or university campus). Thus, we'll also look at how multiple access channels are used in LANs at the end of this section.


We are all familiar with the notion of broadcasting—television has been using it since its invention. But traditional television is a one-way broadcast (that is, one fixed node transmitting to many receiving nodes), while nodes on a computer network broadcast channel can both send and receive. Perhaps a more apt human analogy for a broadcast channel is a cocktail party, where many people gather in a large room (the air providing the broadcast medium) to talk and listen. A second good analogy is something many readers will be familiar with—a classroom—where teacher(s) and student(s) similarly share the same, single, broadcast medium. A central problem in both scenarios is that of determining who gets to talk (that is, transmit into the channel), and when. As humans, we've evolved an elaborate set of protocols for sharing the broadcast channel:

- “Give everyone a chance to speak.”
- “Don’t speak until you are spoken to.”
- “Don’t monopolize the conversation.”
- “Raise your hand if you have a question.”
- “Don’t interrupt when someone is speaking.”
- “Don’t fall asleep when someone is talking.”

Computer networks similarly have protocols—so-called **multiple access protocols**—by which nodes regulate their transmission into the shared broadcast channel. As shown in Figure 5.8, multiple access protocols are needed in a wide variety of network settings, including both wired and wireless access networks, and satellite networks. Although technically each node accesses the broadcast channel through its adapter, in this section we will refer to the *node* as the sending and receiving device. In practice, hundreds or even thousands of nodes can directly communicate over a broadcast channel.

Because all nodes are capable of transmitting frames, more than two nodes can transmit frames at the same time. When this happens, all of the nodes receive multiple frames at the same time; that is, the transmitted frames **collide** at all of the receivers. Typically, when there is a collision, none of the receiving nodes can make any sense of any of the frames that were transmitted; in a sense, the signals of the colliding frames become inextricably tangled together. Thus, all the frames involved in the collision are lost, and the broadcast channel is wasted during the collision interval. Clearly, if many nodes want to transmit frames frequently, many transmissions will result in collisions, and much of the bandwidth of the broadcast channel will be wasted.

In order to ensure that the broadcast channel performs useful work when multiple nodes are active, it is necessary to somehow coordinate the transmissions of the active nodes. This coordination job is the responsibility of the multiple access protocol. Over the past 40 years, thousands of papers and hundreds of PhD dissertations have been written on multiple access protocols; a comprehensive survey of the first 20 years of



**Figure 5.8** ♦ Various multiple access channels

this body of work is [Rom 1990]. Furthermore, active research in multiple access protocols continues due to the continued emergence of new types of links, particularly new wireless links.

Over the years, dozens of multiple access protocols have been implemented in a variety of link-layer technologies. Nevertheless, we can classify just about any multiple access protocol as belonging to one of three categories: **channel partitioning protocols**, **random access protocols**, and **taking-turns protocols**. We'll cover these categories of multiple access protocols in the following three subsections.

Let's conclude this overview by noting that, ideally, a multiple access protocol for a broadcast channel of rate  $R$  bits per second should have the following desirable characteristics:

1. When only one node has data to send, that node has a throughput of  $R$  bps.
2. When  $M$  nodes have data to send, each of these nodes has a throughput of  $R/M$  bps. This need not necessarily imply that each of the  $M$  nodes always